Snowball RL: the regression is real

Same 1,003 validation tasks • native evaluator • 4,012 archived responses replayed exactly

A Full validation over training

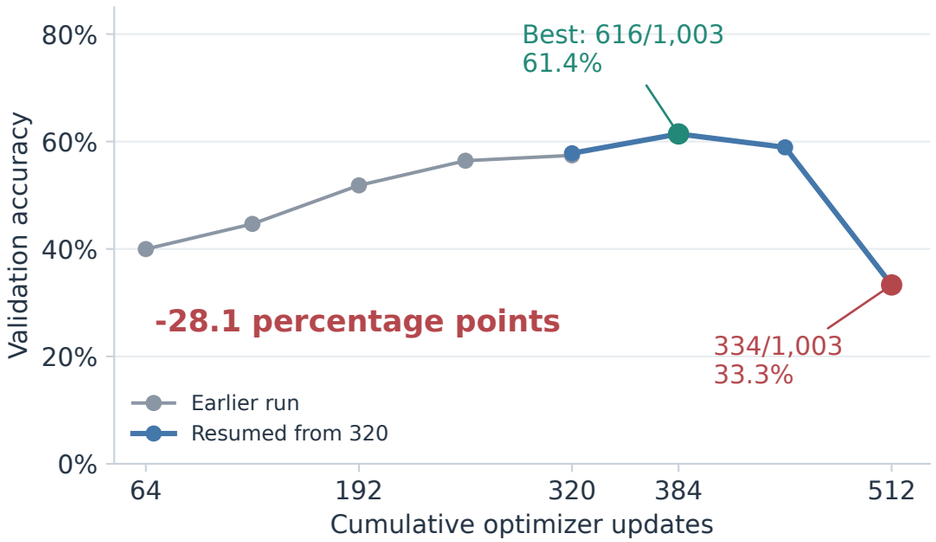

B Answers shorten; truncation falls

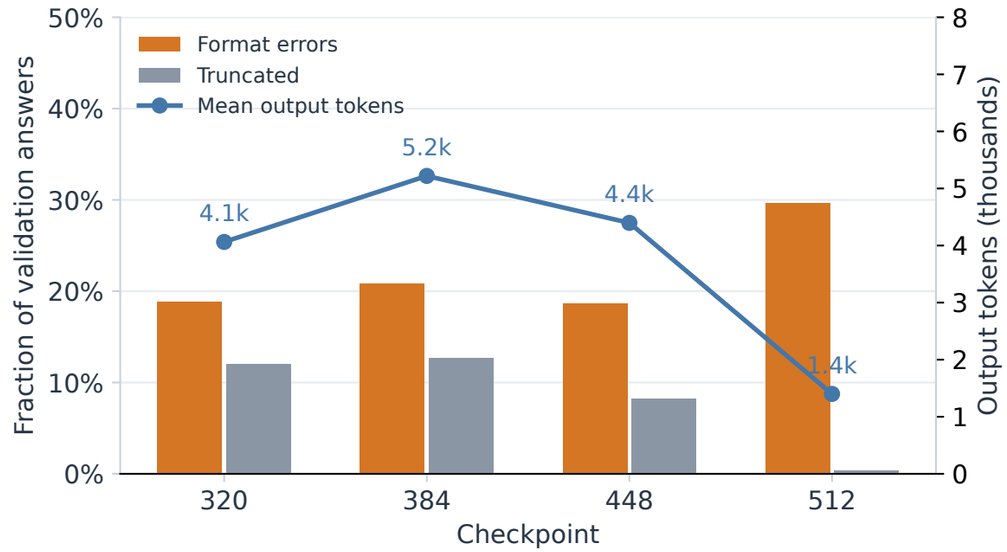

C Geometry and rotamers lose most

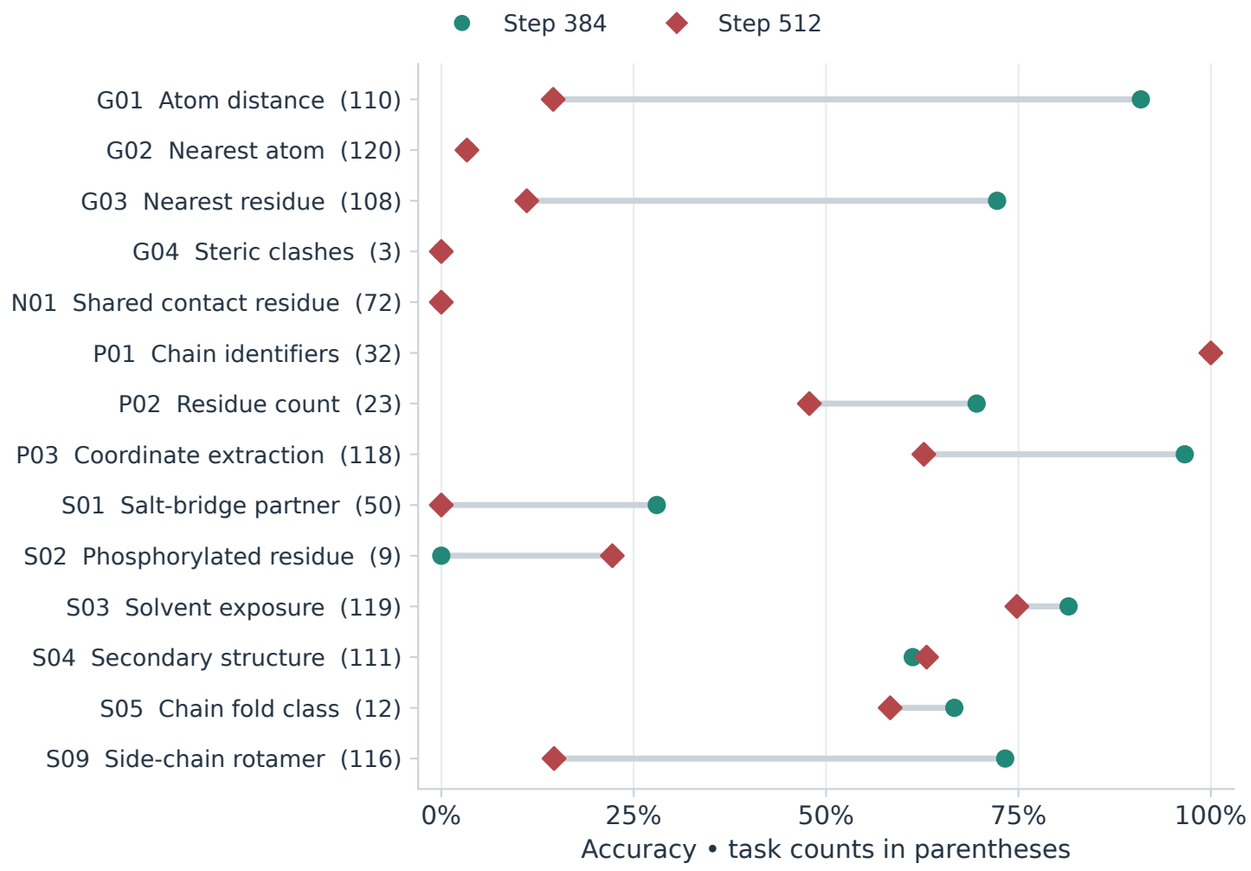

D Most lost answers are valid but wrong

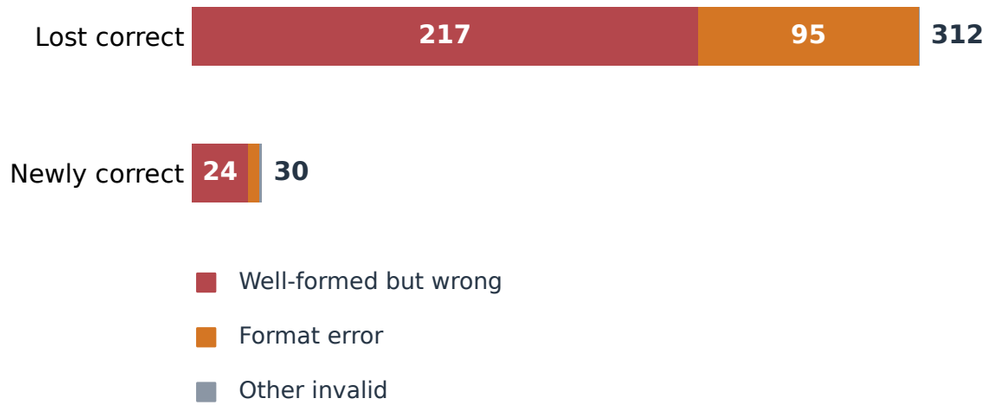

312 losses − 30 gains = 282 fewer correct answers
217/312 losses (70%) are well-formed errors.
All 20 source groups decline.

Step 384 → 512: 95% paired source-group bootstrap interval -33.1 to -24.1 pp. Source-composition uncertainty only.
Native evaluation has a different serving runtime from the original-model/pilot standalone comparison. Do not combine their absolute scores.
Format errors and truncation can overlap. All panels use the frozen native verifier; no new generations.