Does the RL checkpoint improve the original PDBThink benchmark?

Same 40 context-fitting questions · same GB200 inference runtime · checkpoint selected on validation alone

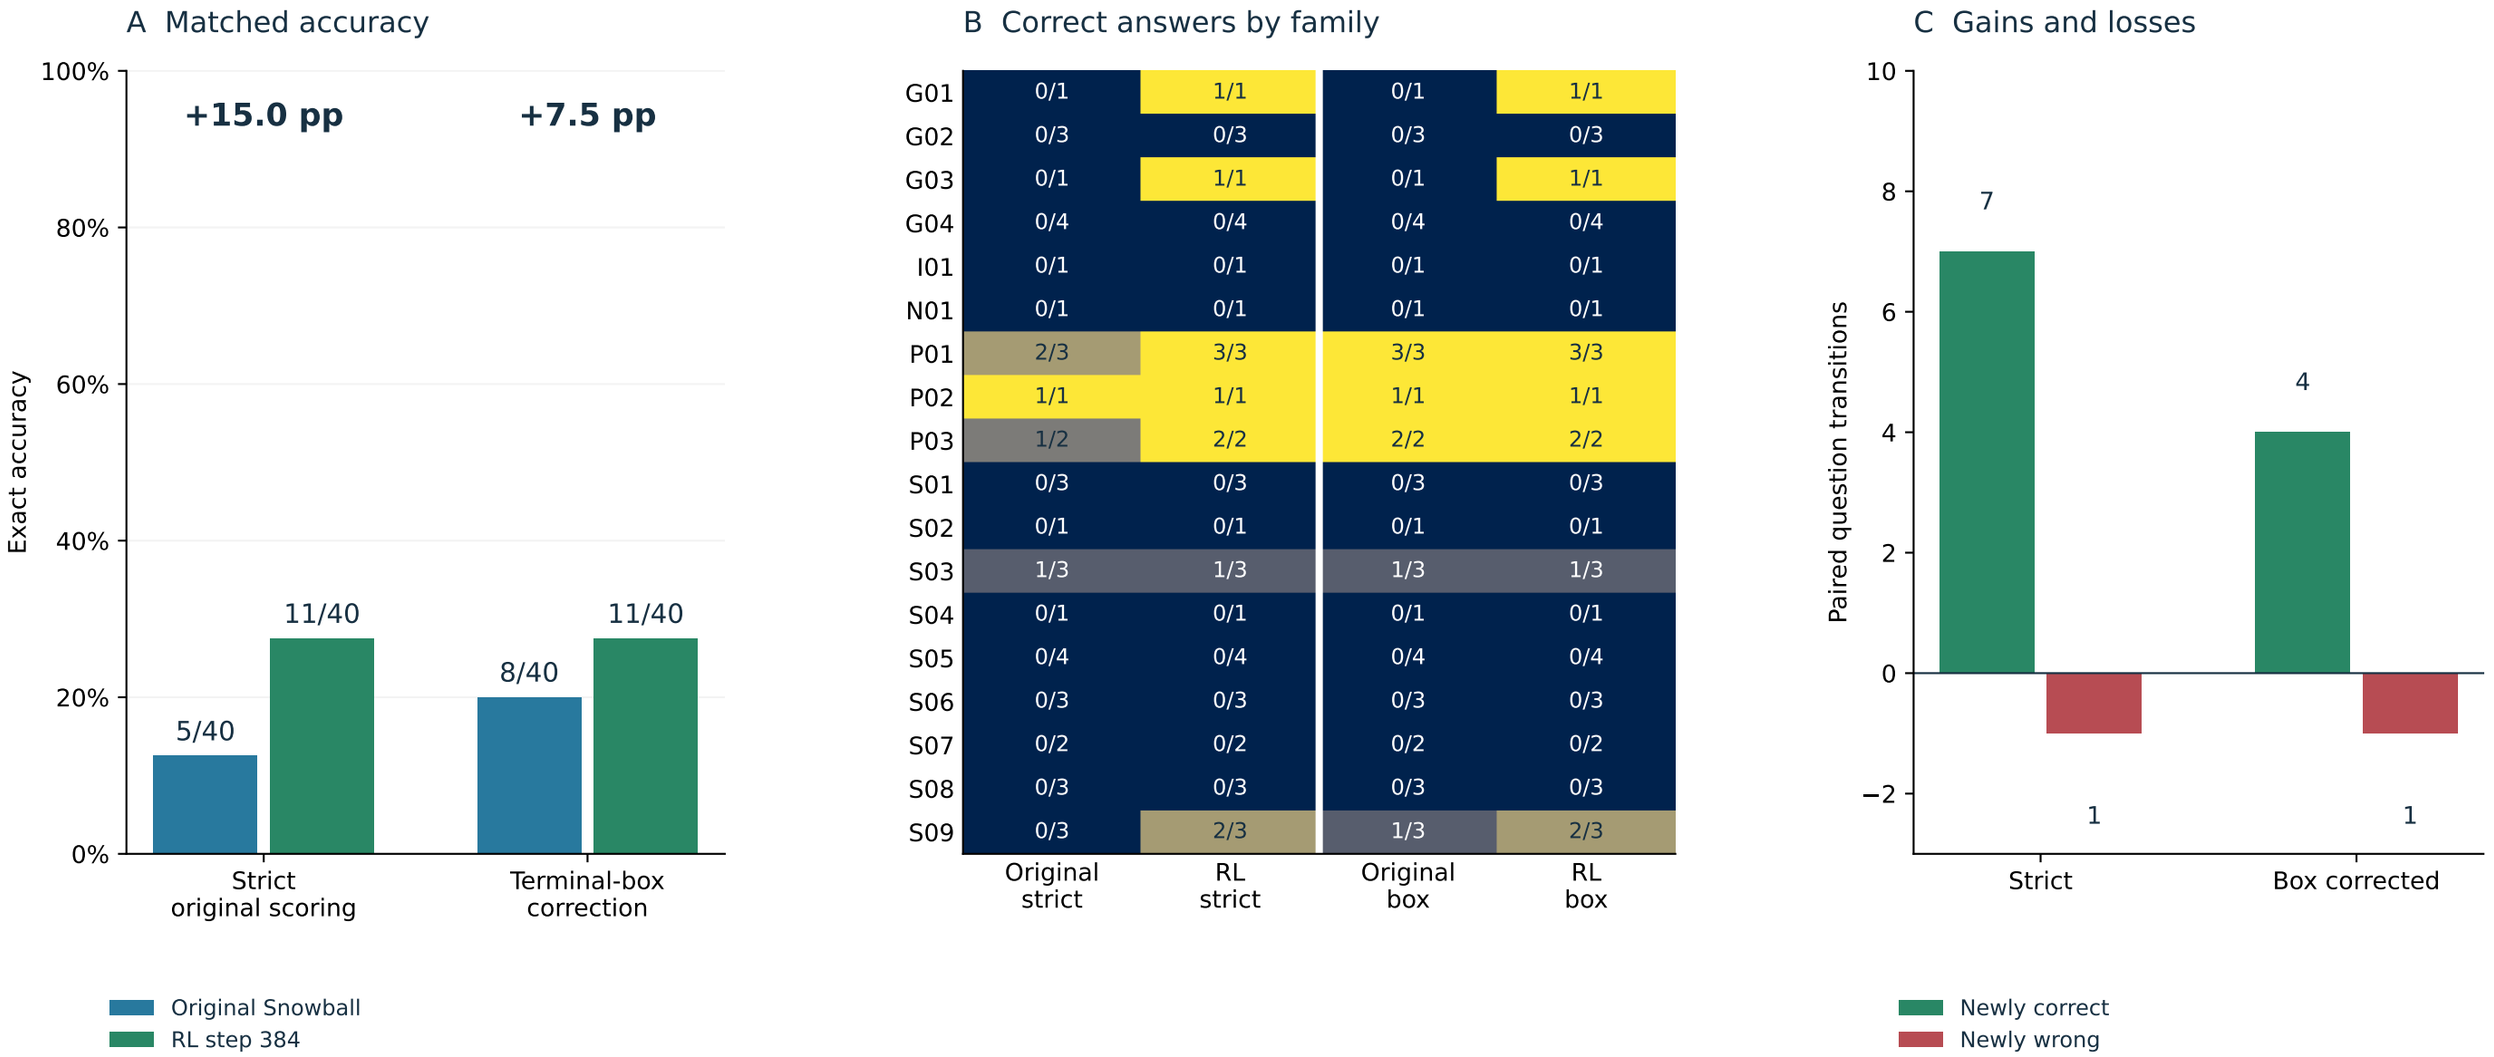

Paired protein-group bootstrap 95% interval for improvement: strict +0.0 to +32.7 pp; box corrected -4.2 to +20.5 pp.

40/88 questions fit; 48 are context-excluded. One response per model/question at temperature 0.6. Intervals exclude training-seed and decoding uncertainty.