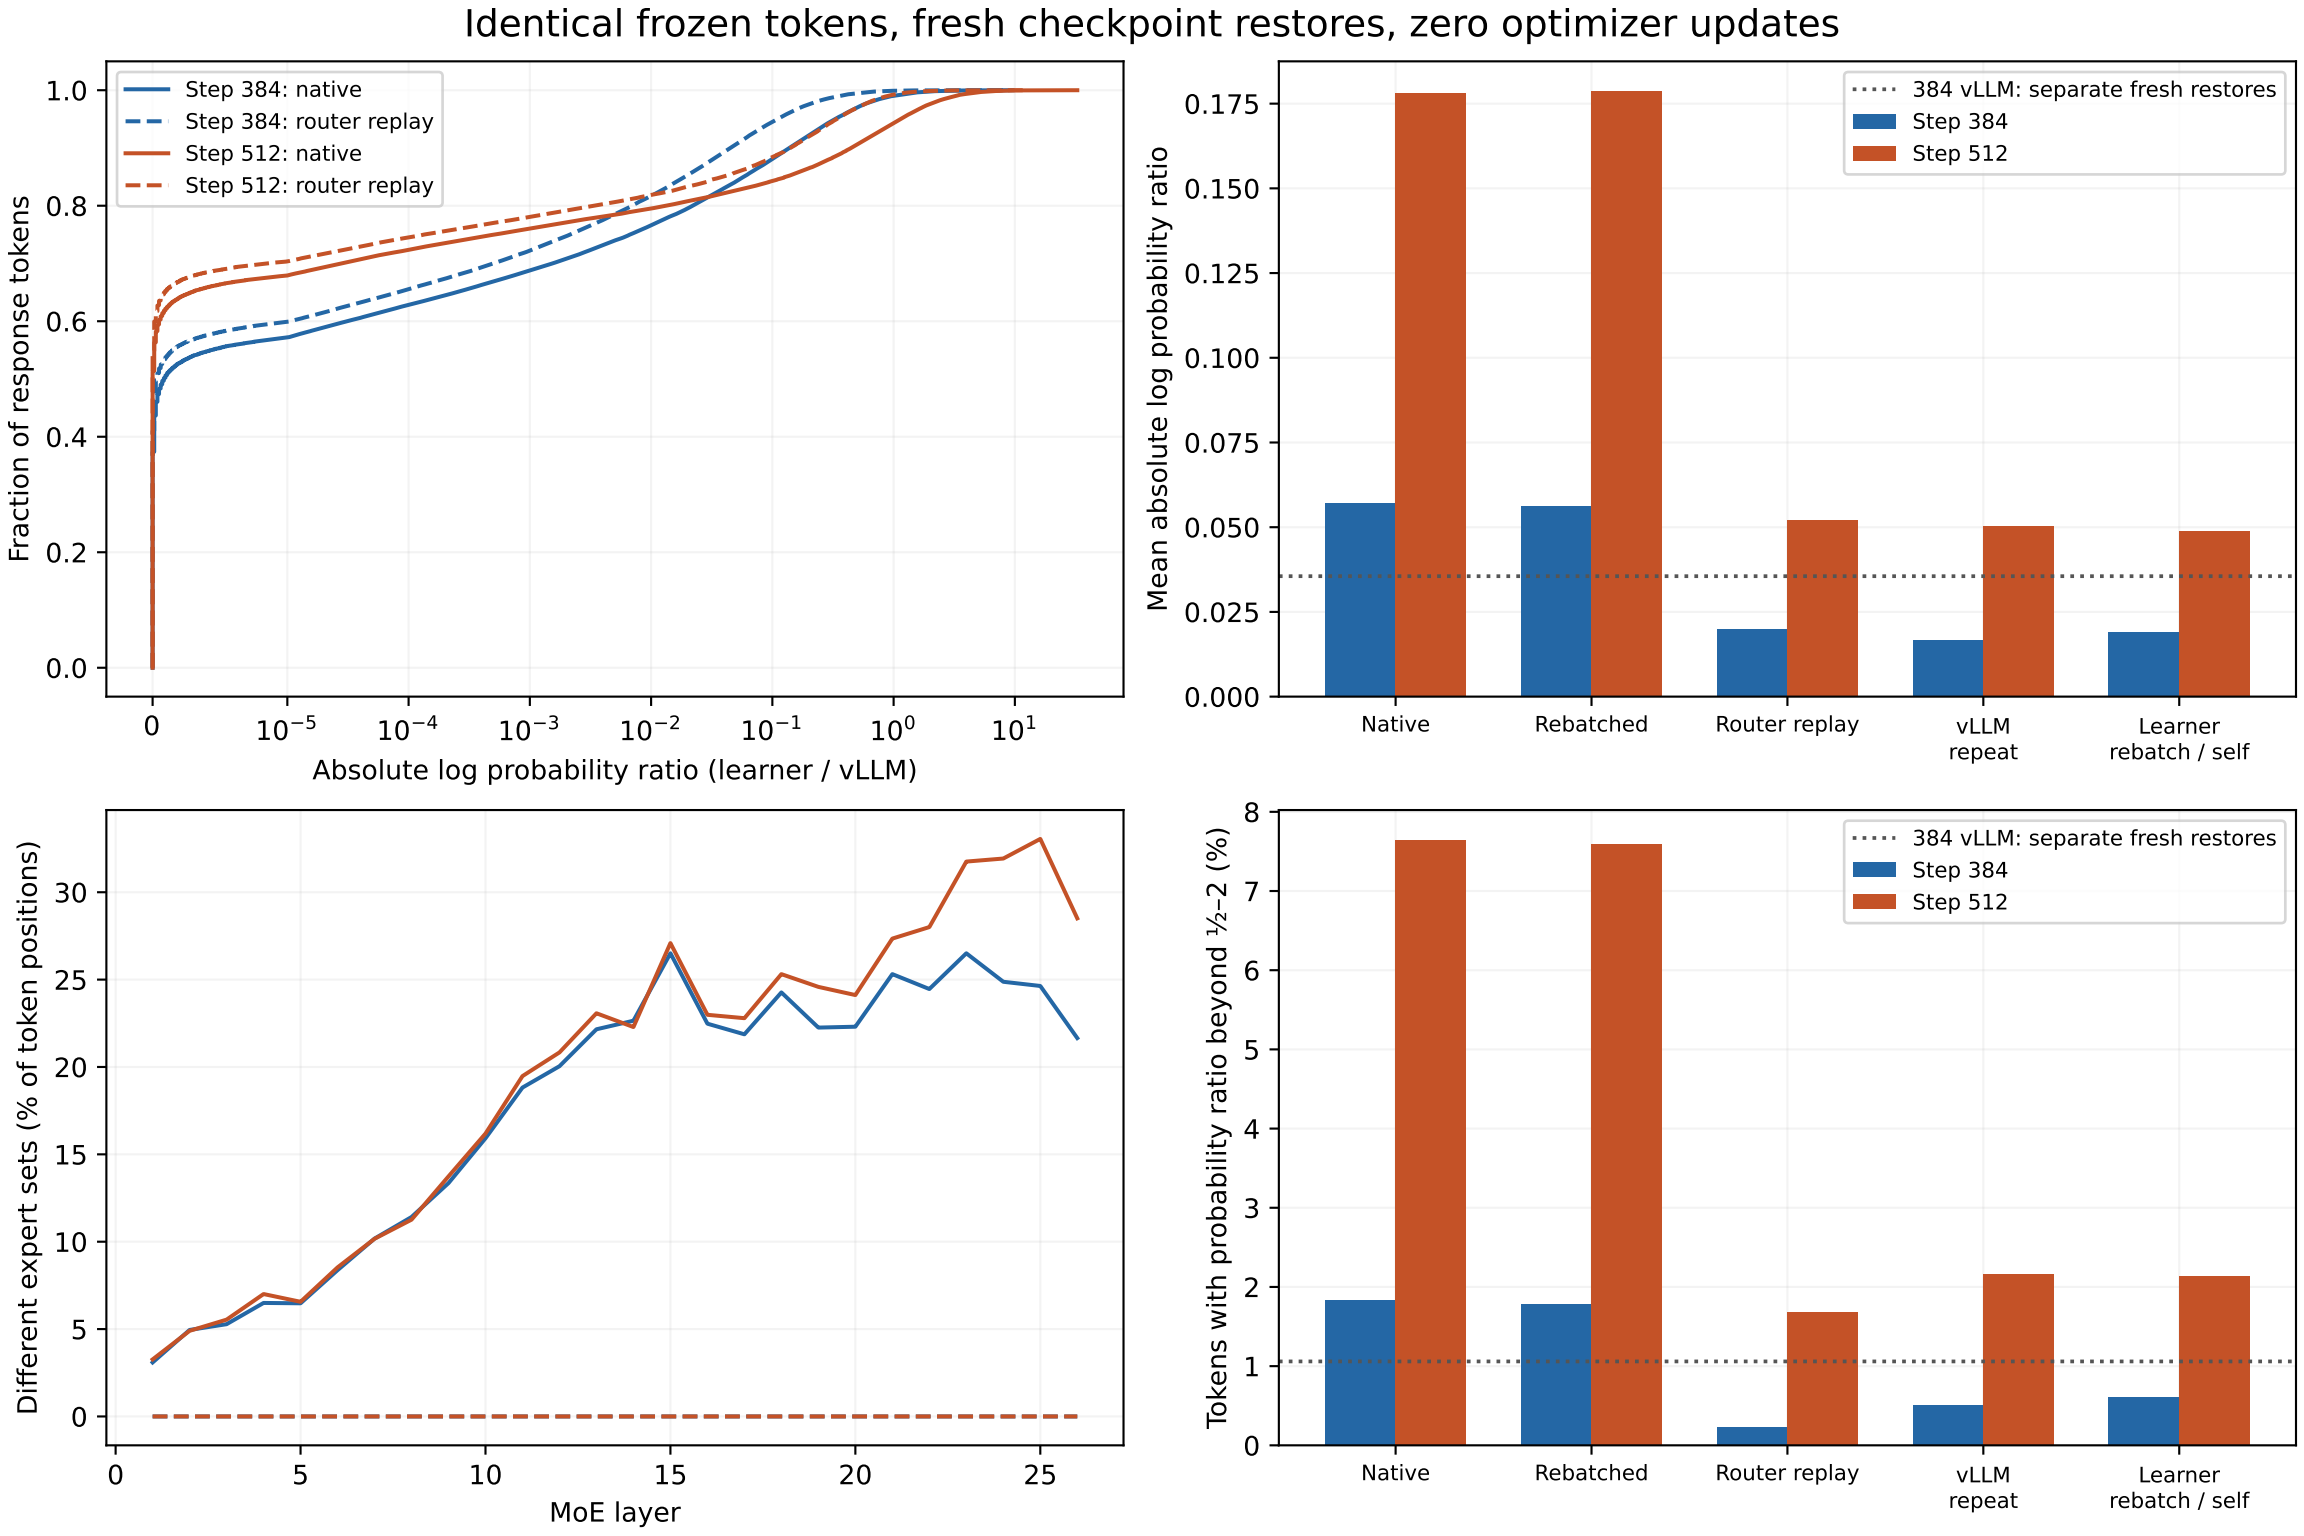
Identical frozen tokens, fresh checkpoint restores, zero optimizer updates
Step 384: native
Step 384: router replay
Step 512: native
Step 512: router replay
Fraction of response tokens
Absolute log probability ratio (learner / vLLM)
384 vLLM: separate fresh restores
Step 384
Step 512
Mean absolute log probability ratio
Native
Rebatched
Router replay
vLLM repeat
Learner rebatch / self
Different expert sets (% of token positions)
MoE layer
Tokens with probability ratio beyond ½–2 (%)